

# CS 480/680

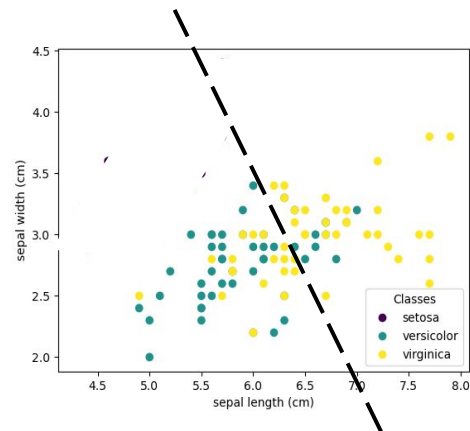
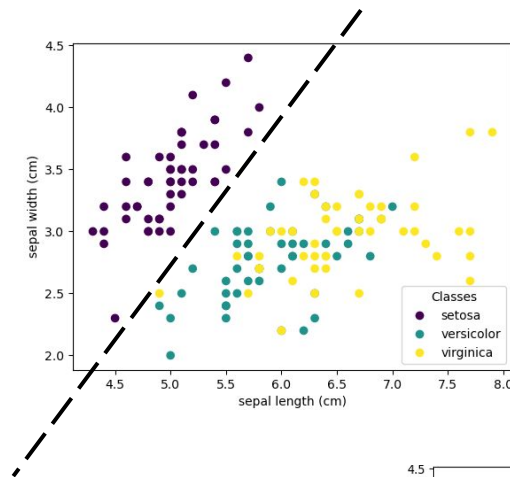
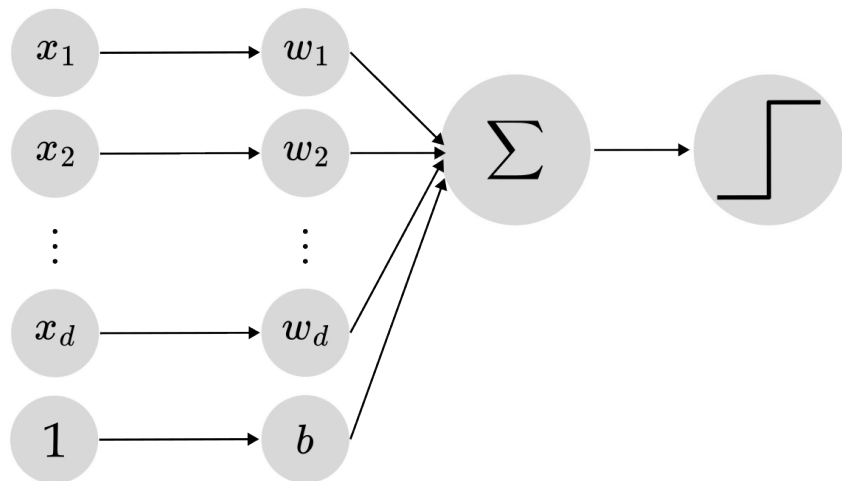
# Introduction to Machine Learning

## Lecture 13

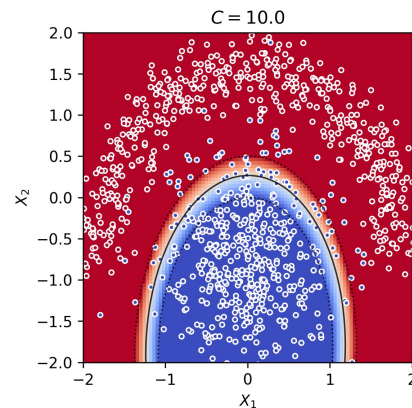
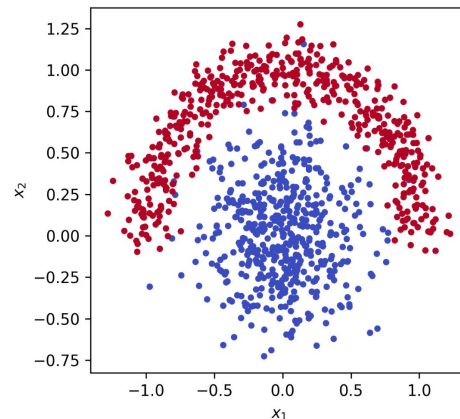
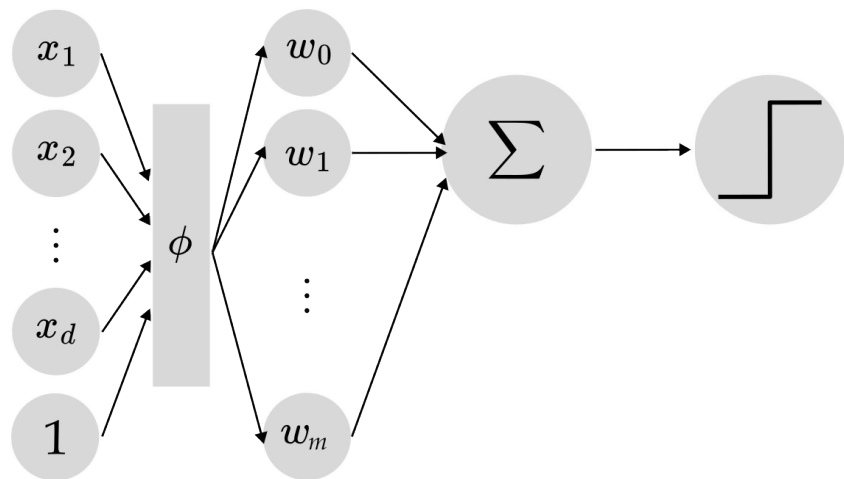
## Multilayer Perceptrons and Deep Learning

Kathryn Simone  
31 October 2024

# Learning functions from linear combinations of inputs

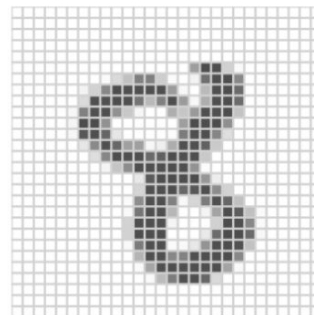
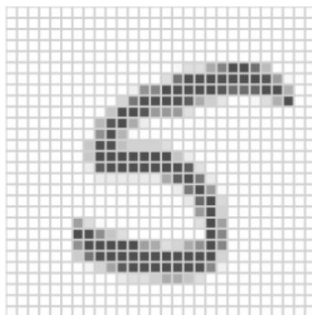
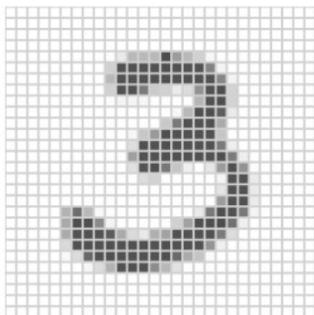


# Learning functions from nonlinear features of inputs

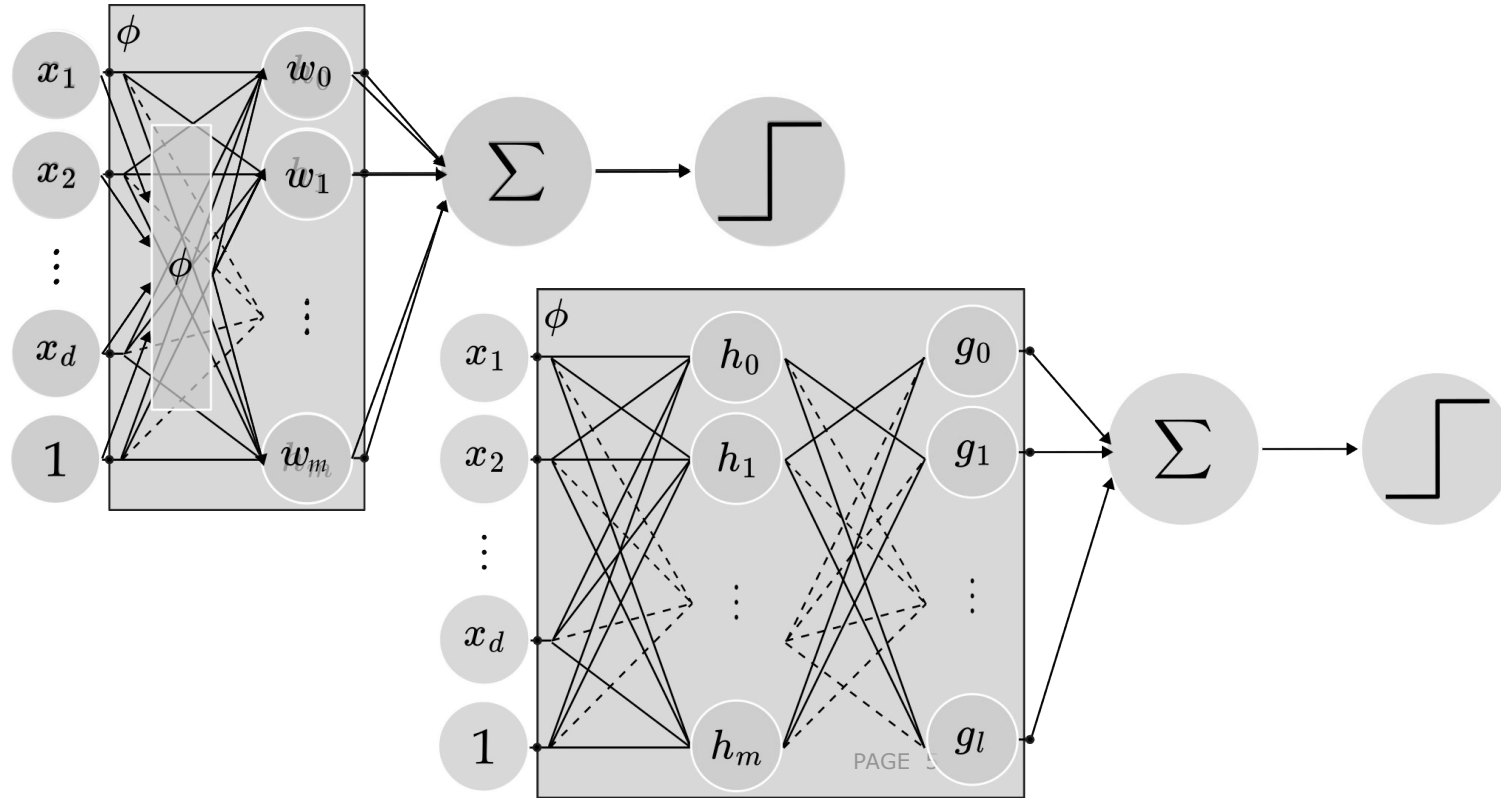


# What if you don't know what the features should be?

0 1 2 3 4 5 6 7 8 9  
0 1 2 3 4 5 6 7 8 9  
0 1 2 3 4 5 6 7 8 9  
0 1 2 3 4 5 6 7 8 9



# Neural Networks: Learn the features of the data from the data



# Key questions

- I. What do we mean by useful features?
- II. How to fit the parameters?
- III. How to prevent overfitting?

# Key questions

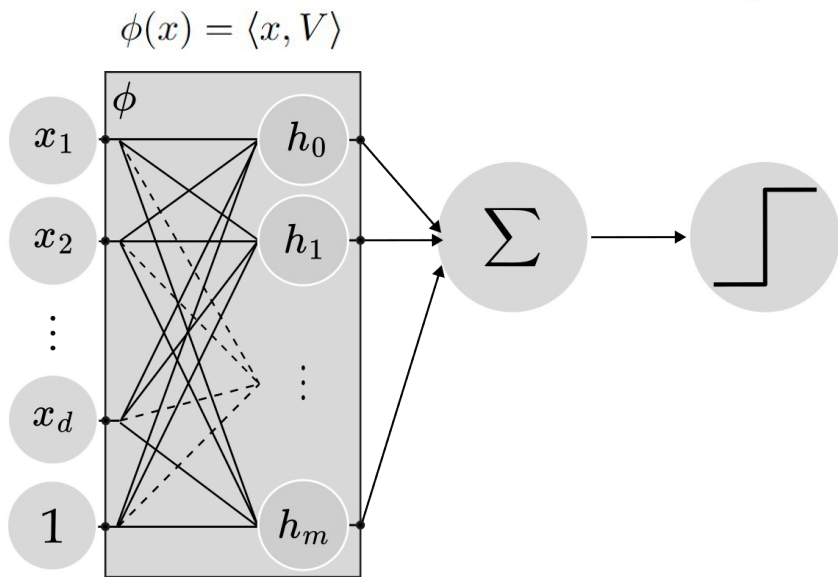
**I. What do we mean by useful features?**

II. How to fit the parameters?

III. How to prevent overfitting?

# How can we represent useful features?

Suppose we create a feature-map  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^m$   $x \in \mathbb{R}^d$  using a *projection matrix*  $V \in \mathbb{R}^{d \times m}$ ,  $V_{ij} \in \mathbb{R}$



$$V = \begin{bmatrix} v_{11} & v_{12} & \dots & v_{1m} \\ v_{21} & v_{22} & \dots & v_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ v_{d1} & v_{d2} & \dots & v_{dm} \end{bmatrix}$$



# How can we approximate features?

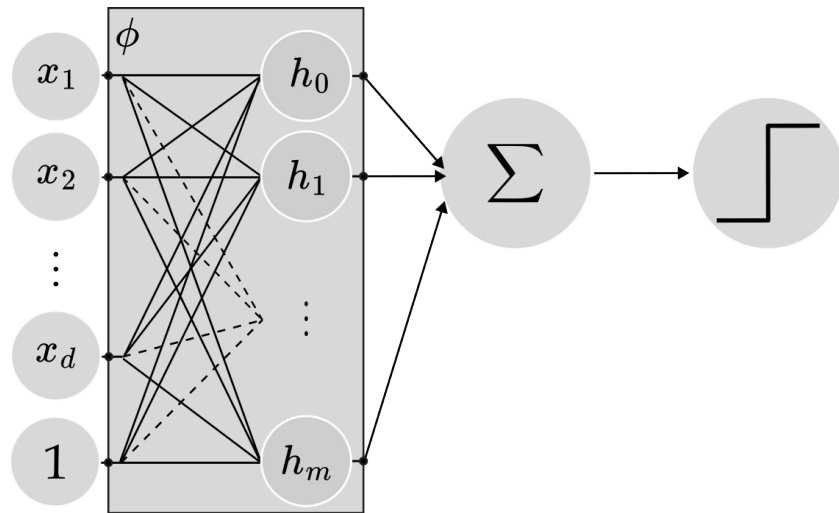
$$\phi(x) = \langle x, V \rangle$$

$$\phi(x) = x^T V$$

$$= \begin{bmatrix} x_1 & x_2 & \dots & x_d \end{bmatrix} \begin{bmatrix} v_{11} & v_{12} & \dots & v_{1m} \\ v_{21} & v_{22} & \dots & v_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ v_{d1} & v_{d2} & \dots & v_{dm} \end{bmatrix}$$

$$= \begin{bmatrix} x_1 v_{11} + x_2 v_{21} + \dots + v_{d1} x_d \\ x_1 v_{12} + x_2 v_{22} + \dots + v_{d2} x_d \\ \vdots \\ x_1 v_{1m} + x_2 v_{2m} + \dots + v_{dm} x_d \end{bmatrix}^T$$

$$= \begin{bmatrix} \phi_1(x) = \sum_{i=1}^d v_{i1} x_i \\ \phi_2(x) = \sum_{i=1}^d v_{i2} x_i \\ \vdots \\ \phi_m(x) = \sum_{i=1}^d v_{im} x_i \end{bmatrix}^T$$

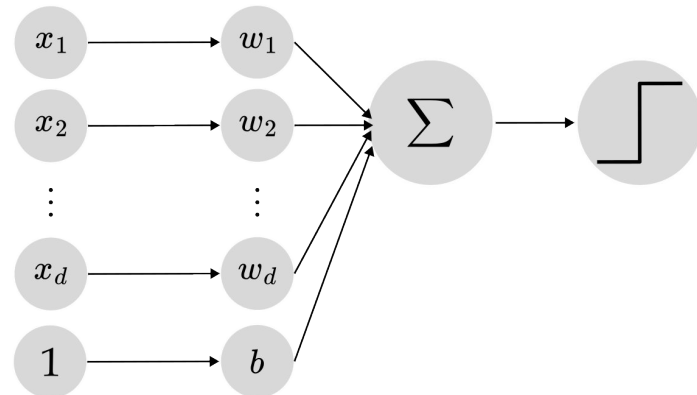
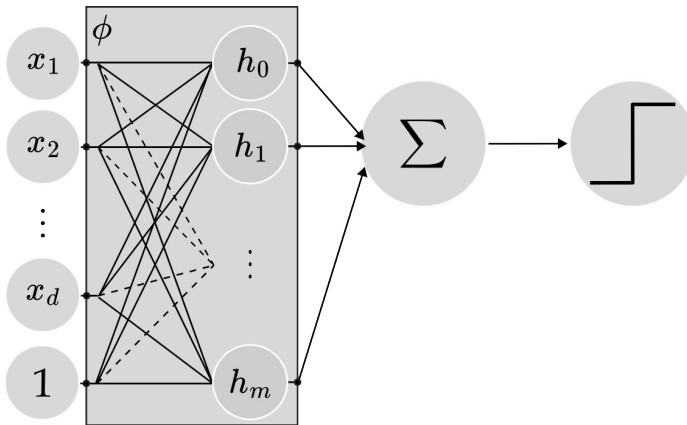


# Linear combinations of input features are not more expressive

$$\begin{aligned}\hat{y} &= \sum_{j=1}^m w_j \phi_j(x) \\ &= \sum_{j=1}^m w_j \sum_{i=1}^d v_{ij} x_i \\ &= \sum_{j=1}^m \sum_{i=1}^d w_j v_{ij} x_i\end{aligned}$$

If we define  $u_i = \sum_{j=1}^m w_j v_{ij}$ , then

$$\hat{y} = \sum_{i=1}^d u_i x_i$$



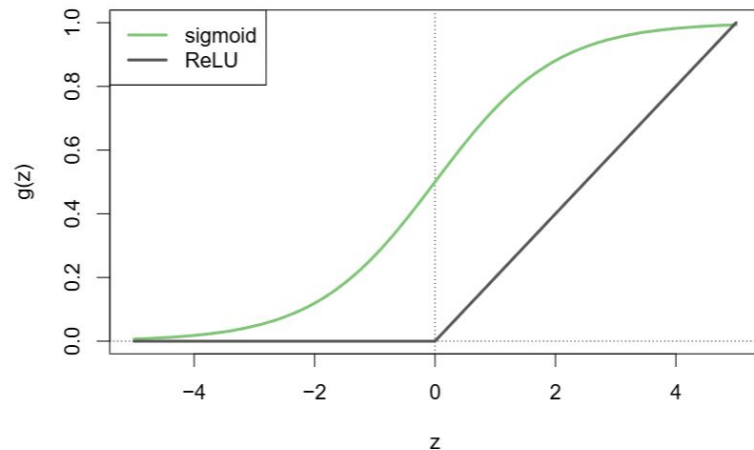
# Solution: Introduce a nonlinearity on hidden units

$$\{\phi_j(x)\} = \left\{ f \left( \sum_{i=1}^d v_{im} x_i \right) \right\} \quad \text{for } j = 1, 2, \dots, m$$

$$\text{Sigmoidal : } \sigma(t) = \frac{1}{1 + e^{-t}}$$

$$\text{Tanh : } \tanh(t) = \frac{e^t - e^{-t}}{e^t + e^{-t}}$$

$$\begin{aligned} \text{Rectified Linear Unit : } \text{ReLU}(t) &= \max(0, t) \\ &= \max(0, t + b) \end{aligned}$$

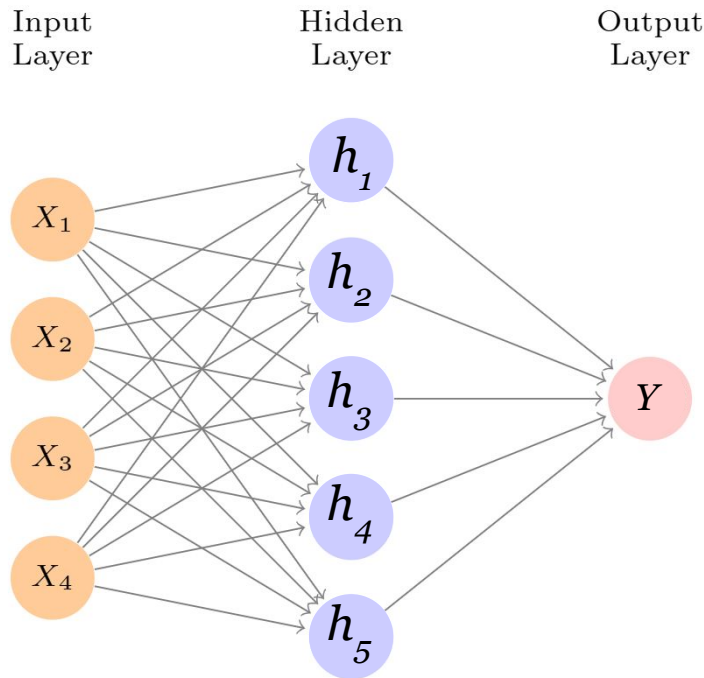


# Universal Approximation Theorem (Hornik 1989, Cybenko 1989)

## Universal Approximation Theorem:

A feedforward network with a linear output layer and at least one hidden layer can approximate any continuous function  $f : \mathbb{R}^d \rightarrow \mathbb{R}^k$  on a closed and bounded subspace of  $\mathbb{R}^d$ , provided:

- ▷ The hidden layer is sufficiently wide
- ▷ The activation function is nonlinear and continuous.



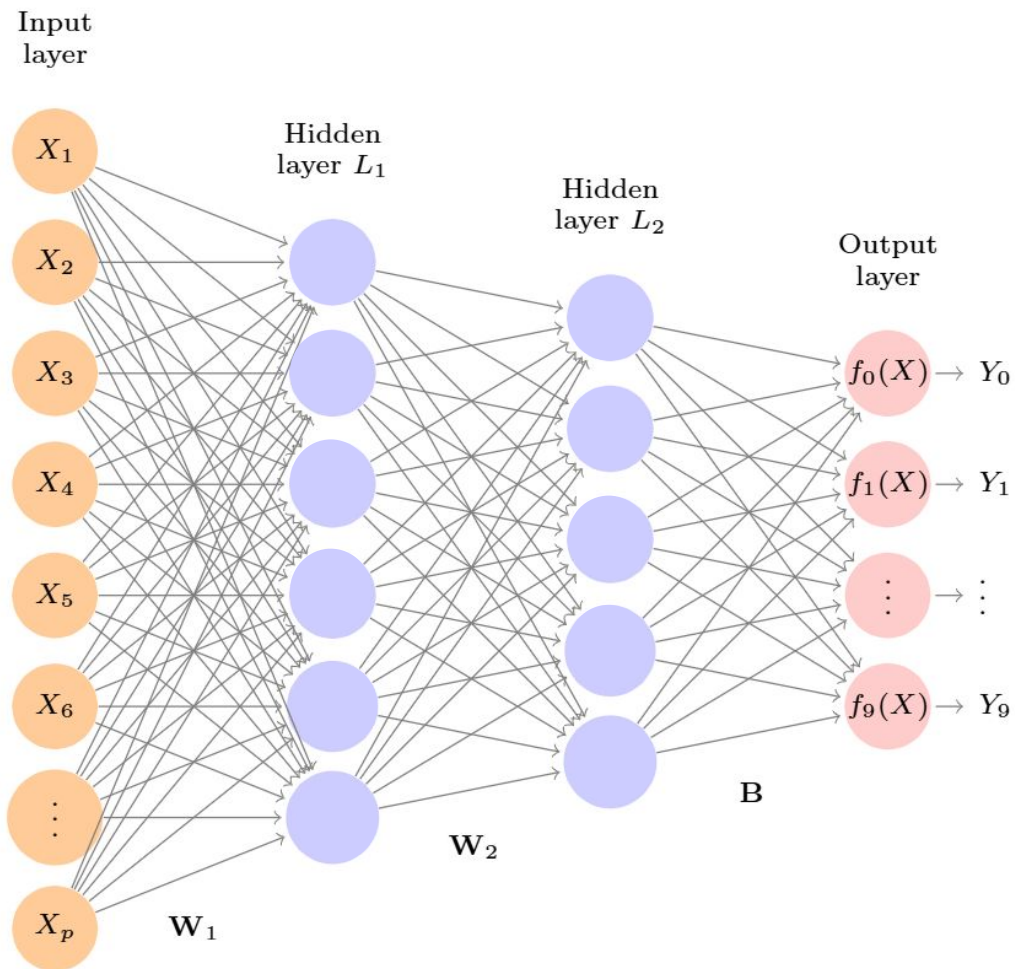
# General Architecture of a Multilayer Perceptron

## Architecture Design:

- Depth: # of hidden layers
- Width: # units per hidden layer

## Functions on output layer

- Typically add a learnable bias
- Sigmoid for classification
- Softmax for multi-class classification
- Could approximate multiple functions at once



# Key questions

I. What do we mean by useful features?

**II. How to fit the parameters?**

III. How to prevent overfitting?

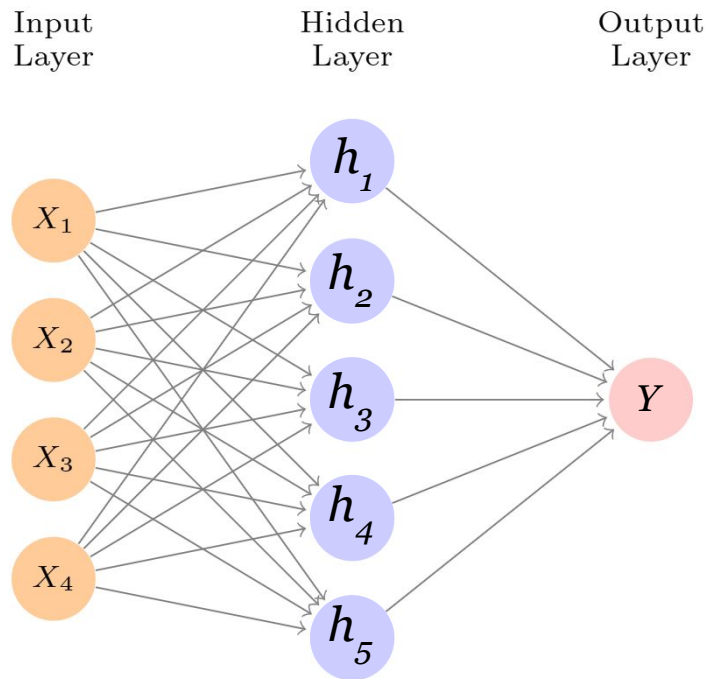
# How to fit the parameters?

Objective:

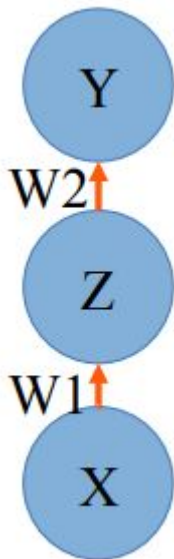
$$\operatorname{argmin}_{\theta} \mathbb{E}[l_{\theta}(x, y)]$$

# trainable parameters:

- Input to Hidden Layer:  
4 x 5 weights/neuron
- Hidden Layer:  
Assuming ReLU, 1 bias/neuron
- Hidden to Output:  
5 weights + 1 bias
- Total: 31 parameters



# Do MLPs have convex objectives?



$$Y = W_1 W_2 X$$

Trained to compute the identity function with squared loss, given a single example pair  $(x_i, y_i) = (1, 1)$

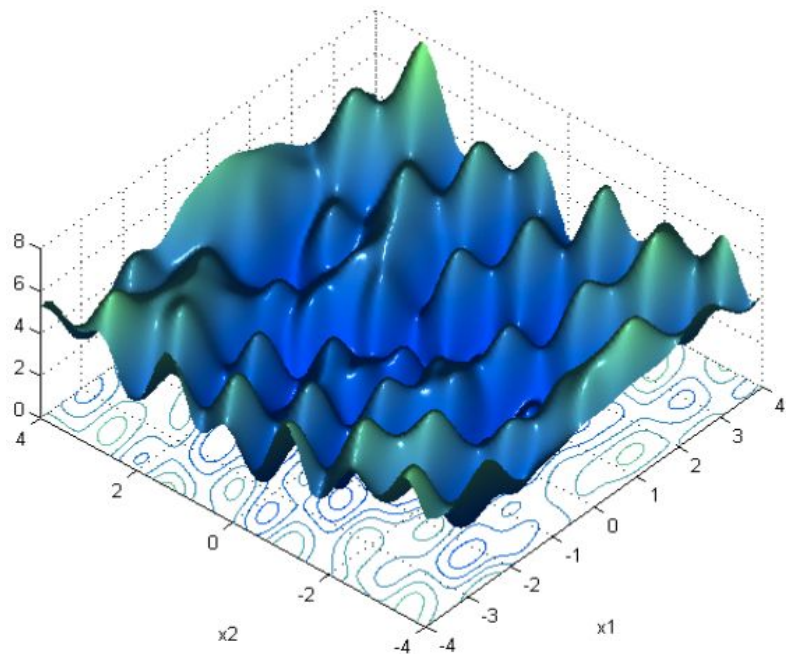
$$L(W) = (1 - W_1 W_2)^2$$

$$\implies W_1 W_2 = 1$$

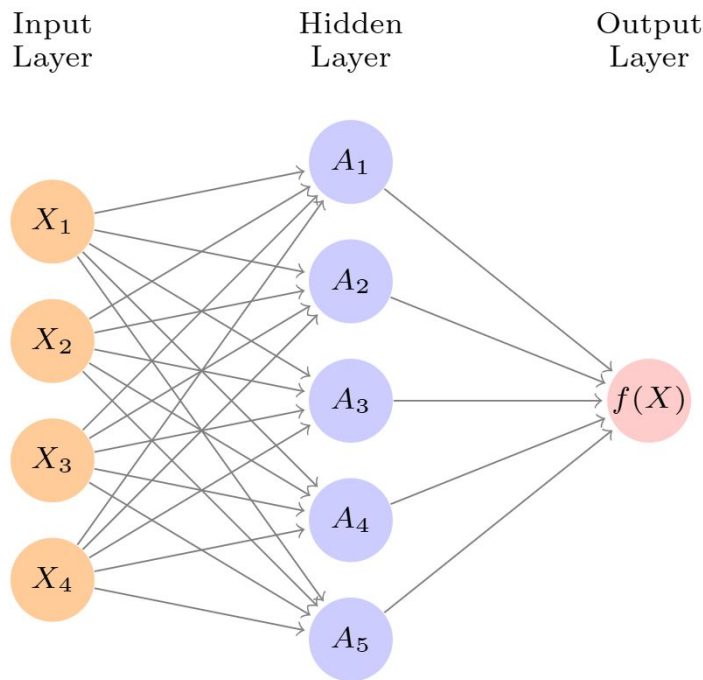
$$\implies W_2 = \frac{1}{W_1}$$



MLPs may have many local minima that are close to global minima



# How to compute the gradient of the loss for an MLP?



# Computing the gradients of an MLP (linear output)

$$L = \frac{1}{2}(\hat{y} - y)^2$$

$$\hat{y} = z$$

$$z = w_1 X_1 + w_2 X_2 + b_{\text{out}}$$

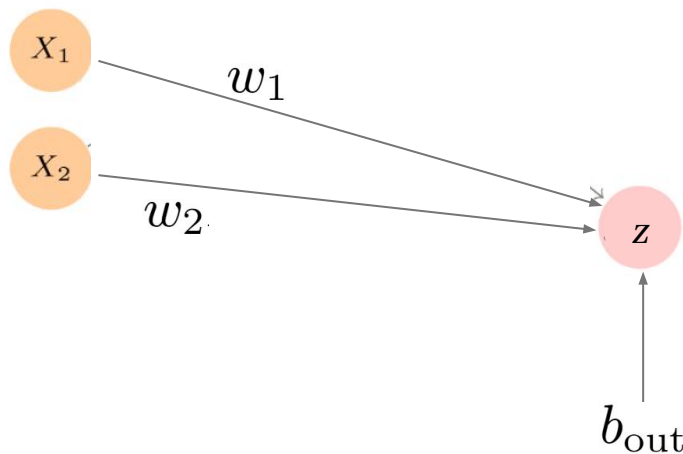
$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial w_1}$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{d}{d\hat{y}} \left( \frac{1}{2}(\hat{y} - y)^2 \right) = (\hat{y} - y)(1) = \hat{y} - y$$

$$\frac{\partial \hat{y}}{\partial z} = \frac{d}{dz}(z) = 1$$

$$\frac{\partial z}{\partial w_1} = \frac{d}{dw_1} \left( w_1 X_1 + w_2 X_2 + b_{\text{out}} \right)$$

$$\frac{\partial L}{\partial w_1} = (\hat{y} - y) \cdot 1 \cdot X_1 = (w_1 X_1 + w_2 X_2 + b_{\text{out}} - y) X_1$$



# Computing the gradients of an MLP (nonlinear output)

$$L = \frac{1}{2}(\hat{y} - y)^2$$

$$\hat{y} = f(z)$$

$$z = w_1 X_1 + w_2 X_2 + b_{\text{out}}$$

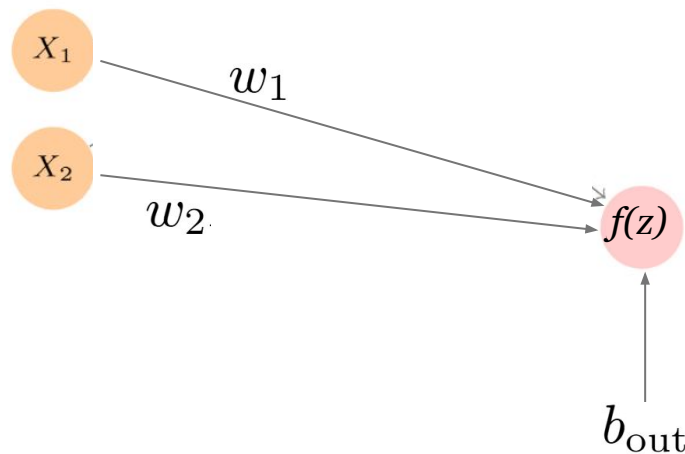
$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial w_1}$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{d}{d\hat{y}} \left( \frac{1}{2}(\hat{y} - y)^2 \right) = (\hat{y} - y)(1) = \hat{y} - y$$

$$\frac{\partial \hat{y}}{\partial z} = \frac{d}{dz} f(z) = f'(z)$$

$$\frac{\partial z}{\partial w_1} = \frac{d}{dw_1} \left( w_1 X_1 + w_2 X_2 + b_{\text{out}} \right) = X_1$$

$$\frac{\partial L}{\partial w_1} = (\hat{y} - y) \cdot f'(z) \cdot X_1$$



# Computing the gradients of an MLP (hidden layer)

$$L = \frac{1}{2}(\hat{y} - y)^2 \quad \hat{y} = f(z)$$

$$z = w_{out}a + b_{out}$$

$$a = h(t)$$

$$t = w_1X_1 + w_2X_2 + b_1$$

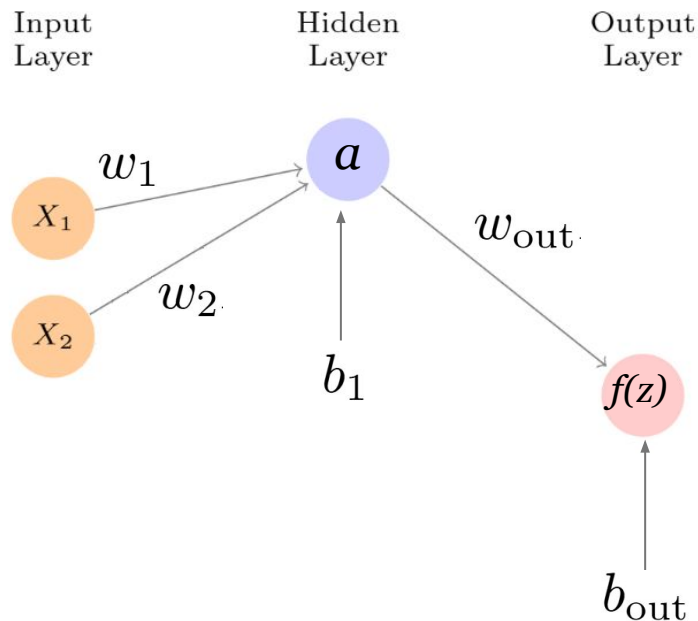
$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial a} \frac{\partial a}{\partial t} \frac{\partial t}{\partial w_1}$$

$$\frac{\partial z}{\partial a} = \frac{d}{da}(w_{out}a + b_{out}) = w_{out}$$

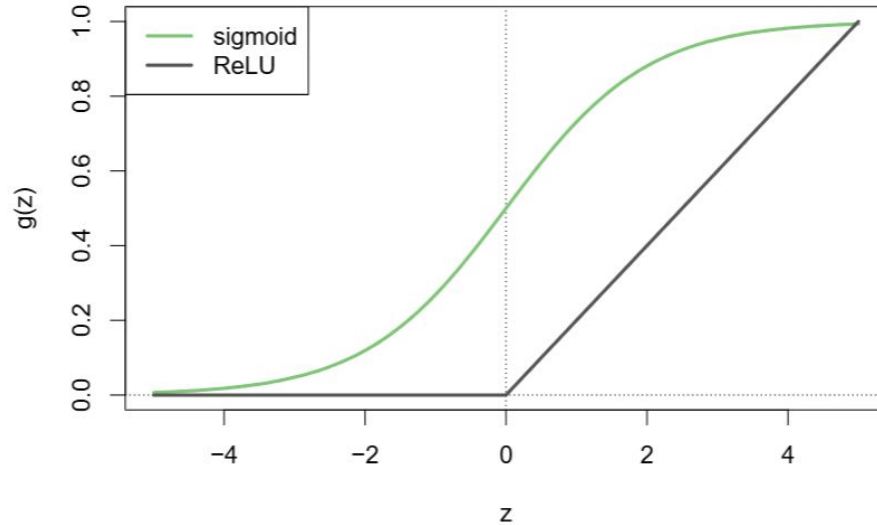
$$\frac{\partial a}{\partial t} = \frac{d}{dt}(h(t)) = h'(t)$$

$$\frac{\partial t}{\partial w_1} = \frac{d}{dw_1}(w_1X_1 + w_2X_2 + b_1) = X_1$$

$$\Rightarrow \frac{\partial L}{\partial w_1} = (\hat{y} - y) \cdot f'(z) \cdot w_{out} \cdot h'(t) \cdot X_1$$



# ReLU activation avoids vanishing gradient



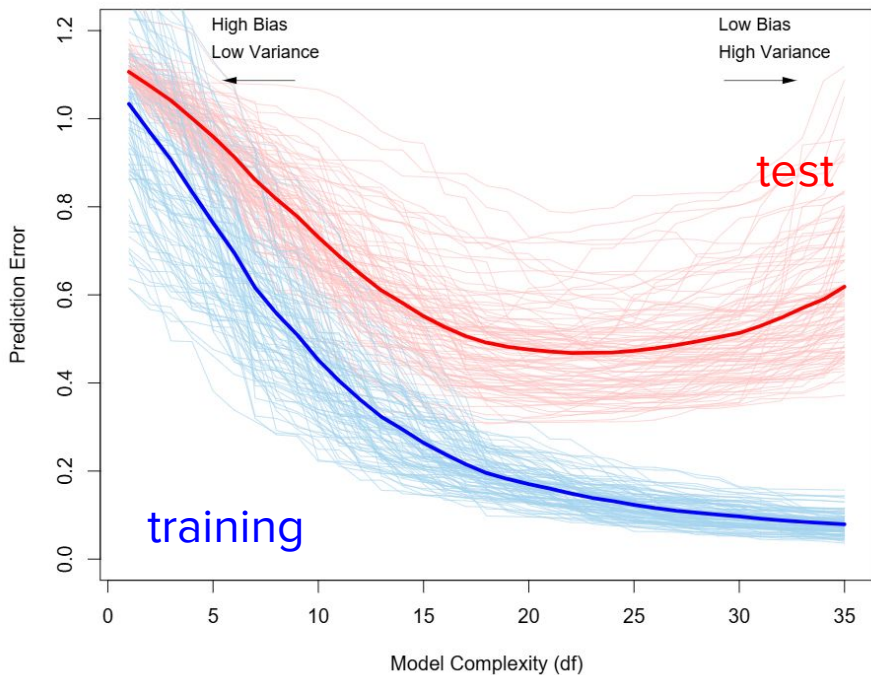
# Key questions

I. What do we mean by useful features?

II. How to fit the parameters?

**III. How to prevent overfitting?**

# Deep neural networks are susceptible to overfitting



$p$  parameters,  $n$  training samples,  $d$  features

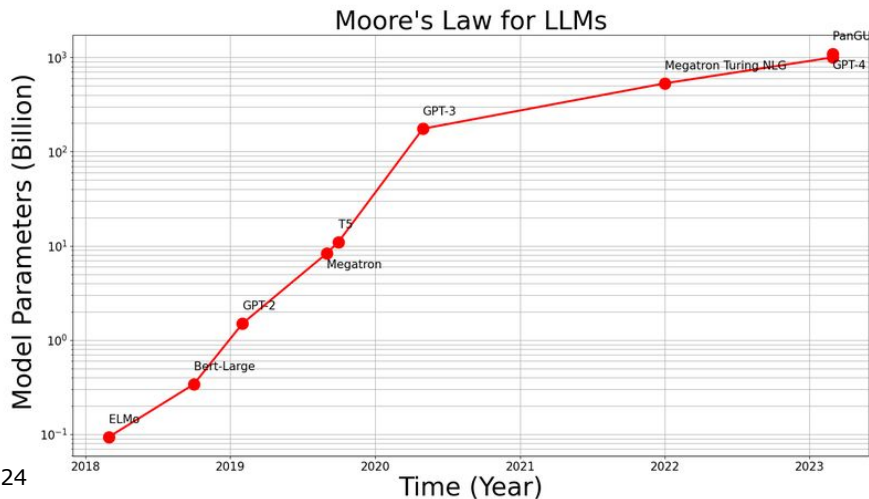
- Classical ML:

$$p \approx d; \ll n$$

- Modern Neural Networks:

$$p \gg nd$$

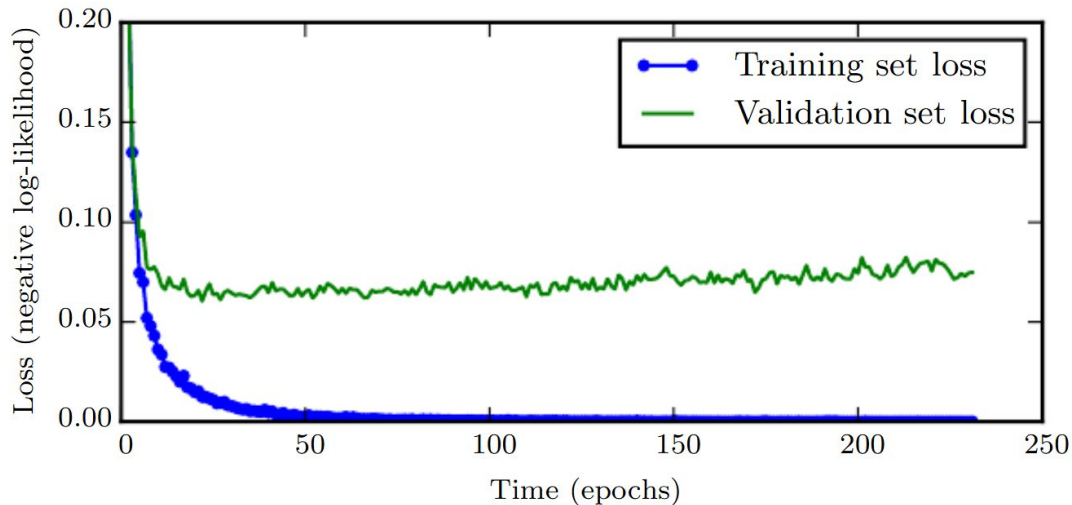
Have sufficient representational capacity to memorize the training dataset





# Regularization Technique #1: Early stopping

- Periodically evaluate the model on a validation set
  - Save a copy of the model if error on the validation set improves
  - Return to the last model checkpoint
- `patience` is number of times to observe worsening validation set error before giving up
- Can think of number of training steps as a hyperparameter



## Regularization Technique #2: Weight decay

$$L = \frac{\lambda}{2} \|w\|^2 + \frac{1}{n} \sum_{i=1}^n L_w(x_i, y_i)$$

$$\nabla_w L = \lambda w + \frac{1}{n} \sum_{i=1}^n \nabla_w L_w(x_i, y_i)$$

$$w_t = w_{t-1} - \eta \left( \lambda w_{t-1} + \frac{1}{n} \sum_{i=1}^n \nabla_{w_{t-1}} L_{w_{t-1}}(x_i, y_i) \right)$$

$$w_t = (1 - \eta\lambda)w_{t-1} - \eta \frac{1}{n} \sum_{i=1}^n \nabla_{w_{t-1}} L_{w_{t-1}}(x_i, y_i)$$

# Regularization Technique #3: Data Augmentation

## Applications

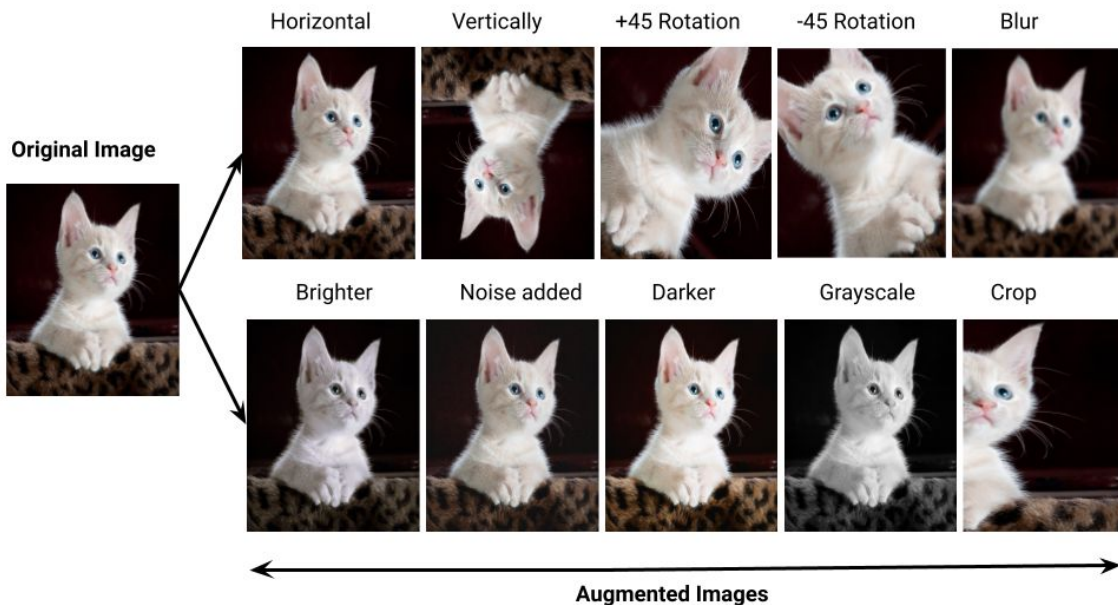
- Classification, object recognition

## Must be used wisely

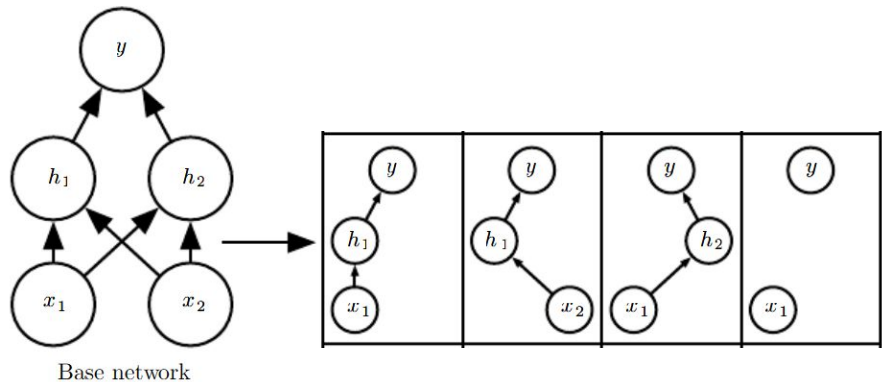
- Certain transformations may change the correct class
- For example, in optical character recognition (OCR):

“6” vs “9” , “b” vs. “d”

Unclear how to apply to tasks like density estimation



# Regularization Technique #4: Dropout (Srivastava 2014)



## Standard drop-down regularization:

- Remove some fraction of  $p$  the nodes in each layer during training
- Applied to each training point separately

$$h' = \begin{cases} 0 & \text{with probability } p \\ \frac{h}{1-p} & \text{otherwise} \end{cases}$$

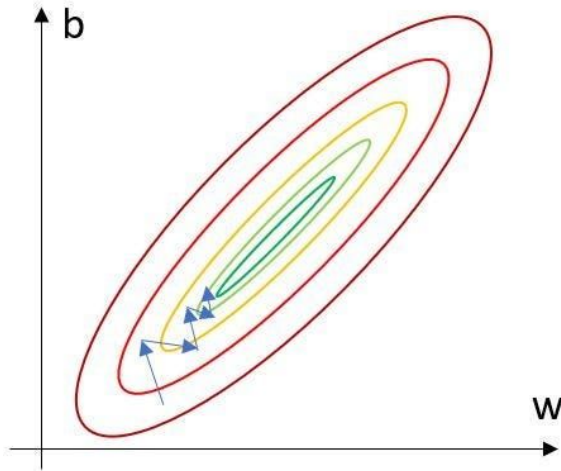
$$E[h'] = h$$

## Advantages

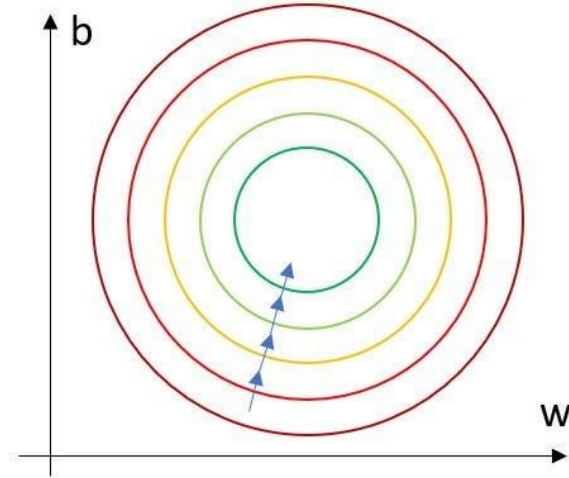
- Computationally cheap
  - Generate  $n$  binary numbers and multiply them by the state
  - $O(n)$  computation/sample/update
  - At inference, only apply the scale
- Can apply different  $p$  to different layers
- Compatible with most architectures, training procedures
- Implemented in most libraries
- Better than weight decay

# Regularization (kindof) Technique #5: Normalization

Unnormalized:



Normalized:



# Regularization (kindof) Technique #5: Normalization

$$\text{BN}(\mathbf{x}) = \gamma \frac{\mathbf{x} - \hat{\mu}_{\mathcal{B}}}{\hat{\sigma}_{\mathcal{B}}} + \beta$$

Where:

$$\hat{\mu}_{\mathcal{B}} = \frac{1}{|\mathcal{B}|} \sum_{\mathbf{x} \in \mathcal{B}} \mathbf{x} \text{ (Batch mean)}$$

$$\hat{\sigma}_{\mathcal{B}}^2 = \frac{1}{|\mathcal{B}|} \sum_{\mathbf{x} \in \mathcal{B}} (\mathbf{x} - \hat{\mu}_{\mathcal{B}})^2 + \epsilon \text{ (Batch variance)}$$

$\gamma, \beta$ : learnable parameters, and

$\epsilon$ : hyperparameter to clip variance estimate.

## Benefits

- Improve optimization by create a consistent scale across features
- Provides a form of regularization by exposing the network to different distributions across batches

## Considerations

- Requires reasonable batch size
- Distributions should not vary significantly between batches (e.g. in reinforcement learning)
- Should be used in combination with other techniques

## Applications to Deep Learning

- Apply to each neuron independently
- Apply to each layer independently

# When is all this effort worth it? A case study

## Baseball Player Salary Dataset

- 263 players
- 19 variables
- 176/87 train/test split
- Goal: Predict player salary

Model	# Parameters	Mean Abs. Error	Test Set $R^2$
Linear Regression	20	254.7	0.56
Lasso	12	252.3	0.51
Neural Network	1345	257.4	0.54

# Attempt neural network after exhausting simpler options

In this case, the Neural Network...

- Achieved similar accuracy
- Required configuring hyperparameters
- Unlike linear methods, parameter values do not suggest a direct interpretation

	Coefficient	Std. error	<i>t</i> -statistic	<i>p</i> -value
Intercept	-226.67	86.26	-2.63	0.0103
Hits	3.06	1.02	3.00	0.0036
Walks	0.181	2.04	0.09	0.9294
CRuns	0.859	0.12	7.09	< 0.0001
PutOuts	0.465	0.13	3.60	0.0005



## Now that we're at the end of the lecture, you should be able to...

- ★ Motivate the need for **learned features**.
- ★ State the **universal approximation theorem** of a multilayer perceptron (MLP) including the conditions (**nonlinear activation, sufficient depth**) and caveats (**potentially infinite width**) to achieve the associated guarantee.
- ★ Use the **chain rule to compute the gradient** with respect to the parameters.
- ★ List and describe theoretical issues (**vanishing gradients, model complexity**) that must be addressed to apply an MLP to real data.
- ★ Describe the **architecture** and **parameterization** of a MLP.
- ★ List, describe, and apply **strategies to reduce overfitting** such as **early stopping, weight decay, data augmentation, drop-out**, as well as **batch/layer normalization**.
- ★ Recommend either a **traditional ML algorithm** or a **neural network** depending on the nature of the problem, the data, and specific goals.