

# CS 480/680

# Introduction to Machine Learning

## Lecture 14

## Convolutional Neural Networks

Kathryn Simone

05 November 2024



UNIVERSITY OF  
**WATERLOO**

FACULTY OF  
MATHEMATICS


# Scaling up Multilayer Perceptrons

- Assume 1 megapixel image:  
 $10^6$  features
- One layer with  $10^3$  units:  
 $10^9$  (1 Billion parameters)
- Challenges:
  - Training cost
  - Require data
  - Risks overfitting



# What properties of images could be exploited?

Locality



# What properties of images could be exploited?

Locality

Spatial Invariance



# Edge detection in the visual system



# The visual system is organized hierarchically




FIGURE 4.8 Response of a single cortical cell to bars presented at various orientations.

# Key questions

- I. How can a network detect low-level features?
- II. How can a network detect higher-level features?

# Convolution is a linear operation over two real-valued functions

$I$ : Input (image)

$K$ : (Convolution) kernel, or filter

$S$ : Output or feature map

$$S(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)$$

**Translation-equivariance:**

If  $g: I'[i, j] \rightarrow I[i-1, j]$ , then  $K \square g(I) = g(K \square I)$



# Can you compute the output of this convolution?

Input

0	1	2
3	4	5
6	7	8

Kernel

0	1
2	3

$*$   $=$  ?

# Convolution kernels *filter* for features in the input

$$K = \begin{bmatrix} 1 & -1 \end{bmatrix}$$

$$K = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$

Input

1	1	0	0	0	0	1	1
1	1	0	0	0	0	1	1
1	1	0	0	0	0	1	1
1	1	0	0	0	0	1	1
1	1	0	0	0	0	1	1
1	1	0	0	0	0	1	1
1	1	0	0	0	0	1	1
1	1	0	0	0	0	1	1


Output, Filter 1

0	1	0	0	0	-1	0
0	1	0	0	0	-1	0
0	1	0	0	0	-1	0
0	1	0	0	0	-1	0
0	1	0	0	0	-1	0
0	1	0	0	0	-1	0
0	1	0	0	0	-1	0
0	1	0	0	0	-1	0


Output, Filter 2

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

# Padding prevents loss of input pixels




# A filter is realized in a CNN through sparse weights



## “Fully-Connected” Neural Network


- Every output interacts with every input



## Convolutional Neural Network


- Only three outputs are affected by a given input

# A filter is realized in a CNN through sparse weights



## “Fully-Connected” Neural Network


- Every output interacts with every input



## Convolutional Neural Network


- Only three inputs affect an output

# Parameter sharing achieves translation-equivariance



## “Fully-Connected” Neural Network


- Interactions have distinct parameters




## Convolutional Neural Network

- Parameters are shared

# The network must be able to detect a variety of patterns

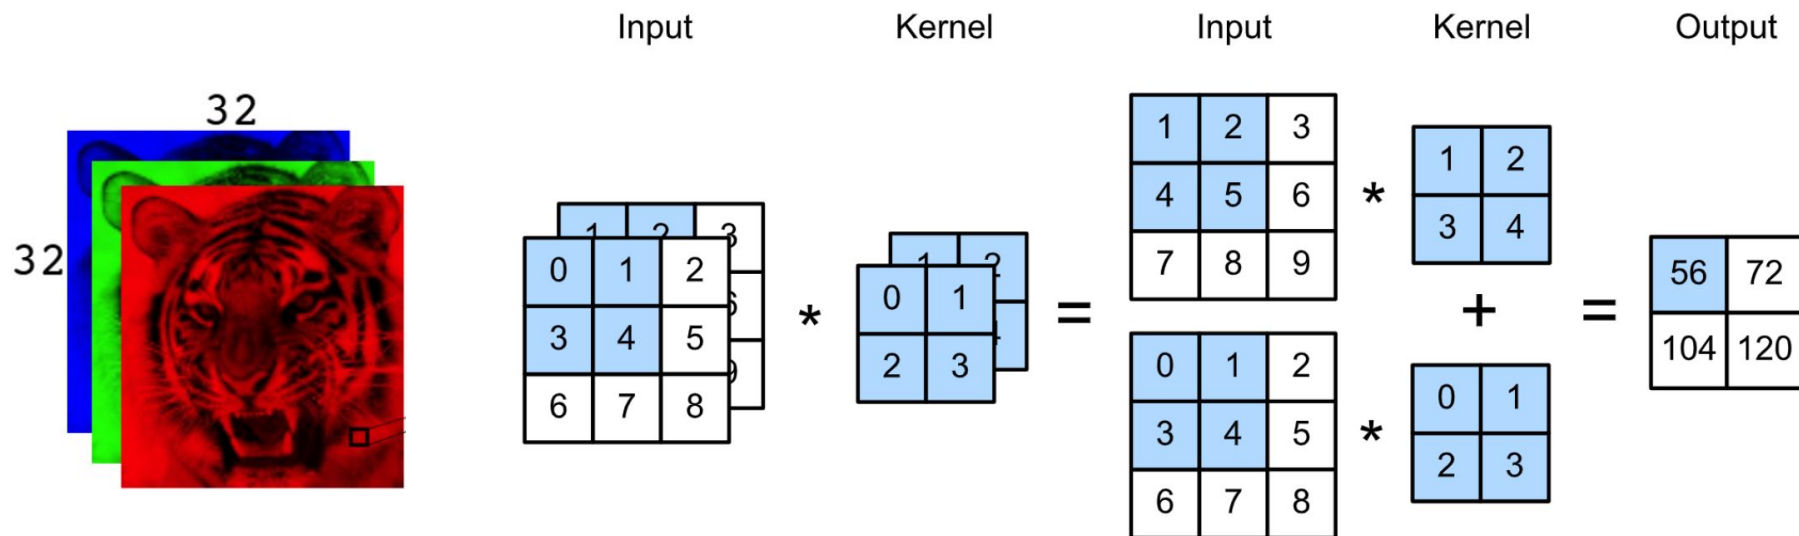


A convolutional *stage* generates several feature maps






# Convolution kernels with input multiple channels



## Stride: Trade off spatial resolution for output channel depth



# Dimensionality and parameters of a convolution stage

Input:

$$W_{IN} \times H_{IN} \times D_{IN}$$

Hyperparameters:

Number of filters  $N_k$

Kernel Size  $F \times F \times D_{IN}$

Padding  $P$

Stride  $S$

Output:

$$W_{OUT} \times H_{OUT} \times D_{OUT}$$

$$W_{OUT} = (W_{IN} + 2P - F) / S + 1$$


$$H_{OUT} = (H_{IN} + 2P - F) / S + 1$$

$$D_{OUT} = N_k$$

# Trainable parameters:

$$\begin{aligned} & N_k \times F \times F \times D_{IN} \text{ weights} \\ + & N_k \text{ biases} \end{aligned}$$

# Convolution is translation equivariant but not invariant




Cat



Cat

# Convolution is not invariant to other transformations



# Key questions

- I. How can a network detect low-level features?
- II. How can a network detect higher-level features?**

# Recall simple and complex cells






FIGURE 4.8 Response of a single cortical cell to bars presented at various orientations.

# Fukushima's Neocognitron (1980) and pooling








# *Pooling* over spatial regions confers translation invariance




# Pooling over feature maps yields latent object representations



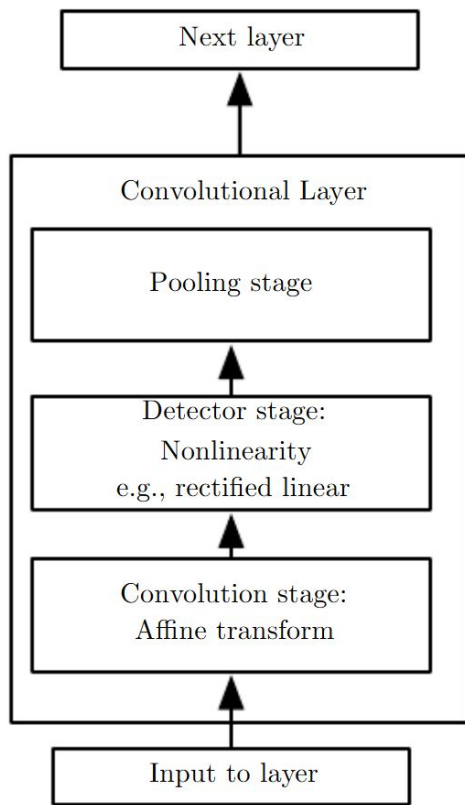
# Invariance to small translation confers invariance to distortion



# From Neocognitron to CNN (LeCun's LeNet, 1989)





## A *convolutional layer* may apply multiple nonlinearities




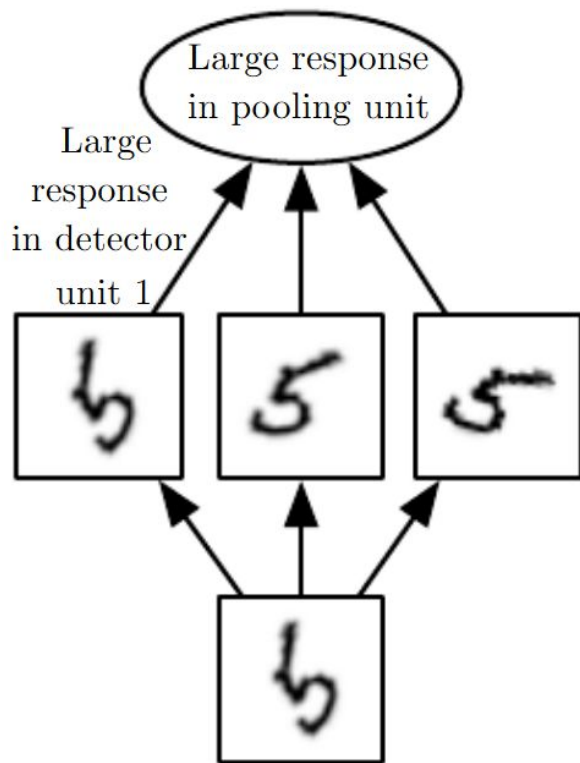
### Pooling stage

- Summarizes nearby outputs
- Average pooling, L2 norm, Max pooling


Pooling across spatial regions induces some translation *invariance*



## Pooling stages summarize their inputs

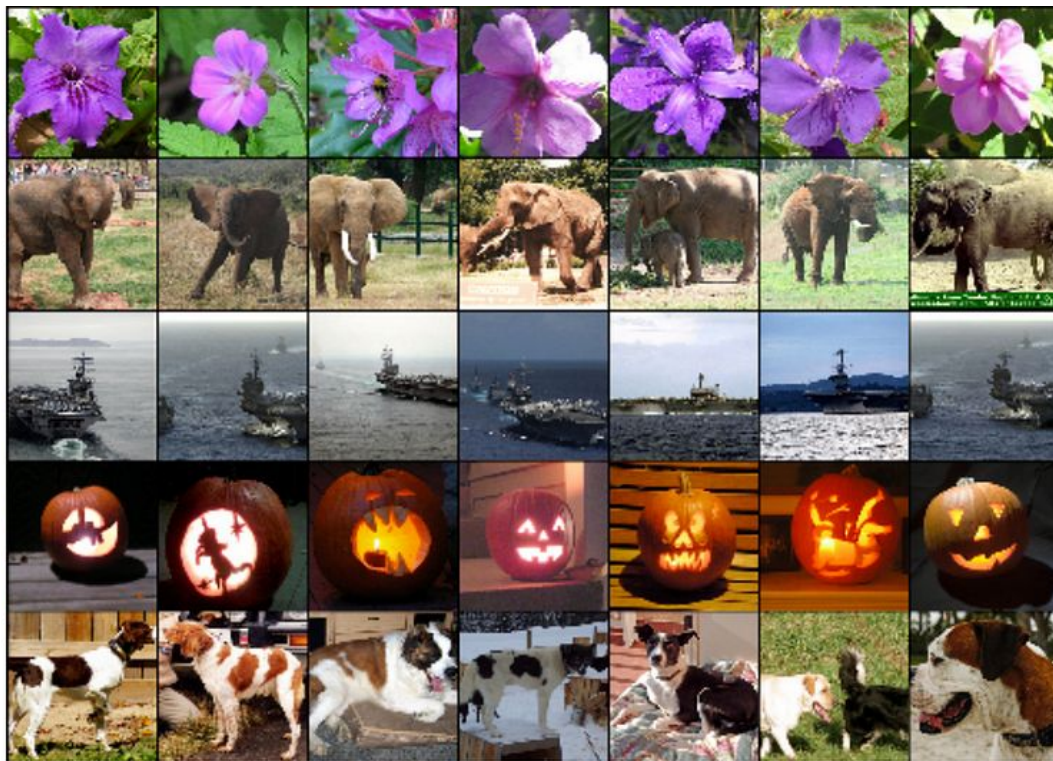


Deeper layers have indirect interactions with most of the input







# CNNs learn a feature space where inner products are meaningful



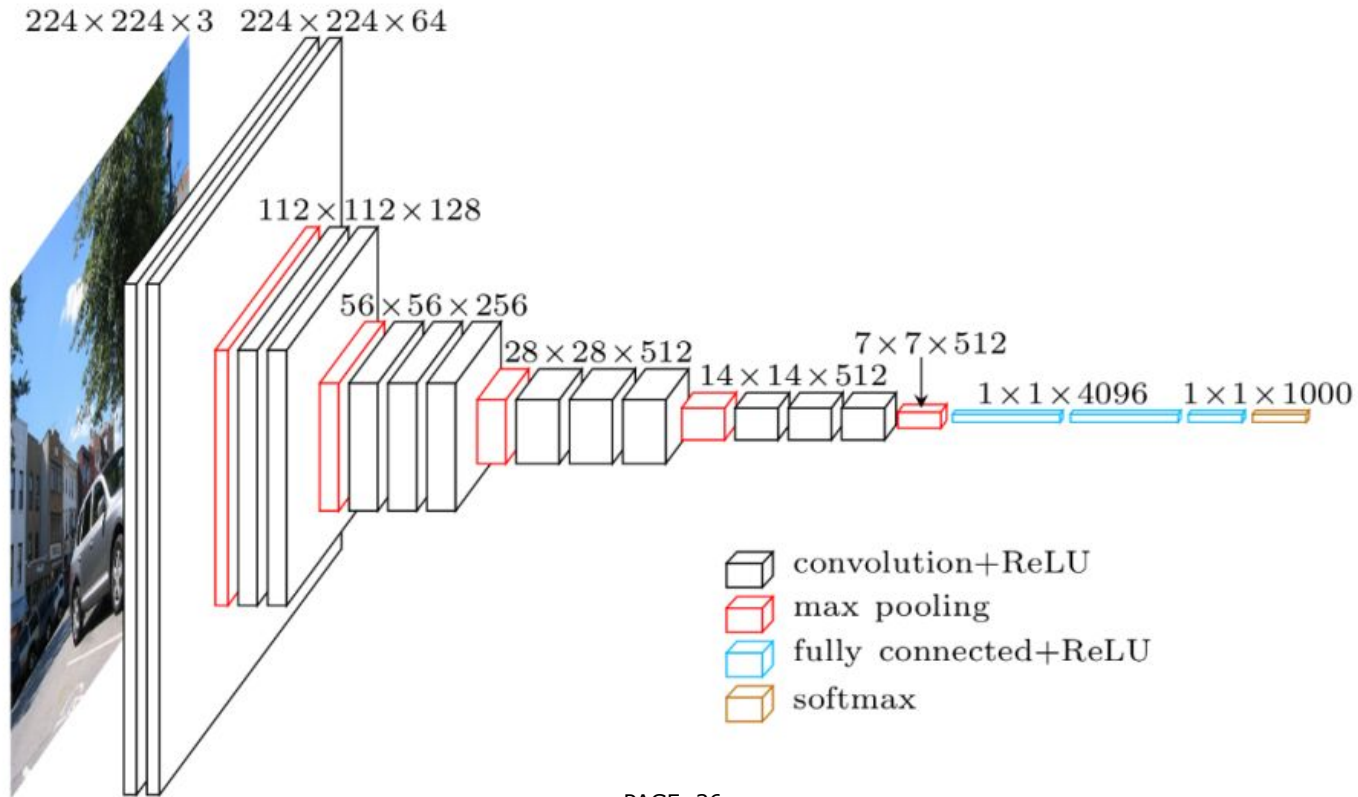
# CNNs learn a feature space where inner products are meaningful



# Interpreting the architecture of a CNN



# Interpreting the architecture of a CNN



## Now that we're at the end of the lecture, you should be able to...

- ★ Compute the grayscale value of an image pixel through **cross-correlation** of a filter and an image.
- ★ Give examples of **low-level filters** used in a CNN for image processing and define their effect with reference to **convolution**.
- ★ Define and explain the role of **parameter sharing** and **local connections** in CNNs.
- ★ Construct a convolutional layer using **convolution**, **nonlinearity** and **pooling** operations.
- ★ Describe the information processing occurring at each layer of a **Convolutional Neural Network (CNNs)**, given a schematic of its architecture.
- ★ Sketch the **architecture of a CNN** given its verbal description.
- ★ Determine the number of trainable parameters in a convolutional layer given the hyperparameters to set **# of filters, kernel size, padding, stride**.