

CS 480/680

Introduction to Machine Learning

Lecture 18

Variational Autoencoders and Normalizing Flows

Deep Generative Models Part I

Kathryn Simone

19 November 2024

Two significant and distinct goals in machine learning

Discriminative Model:

Learn a predictor given the observations.

Examples:

Perceptron, Support Vector Machines
Decision Trees, MLPs, CNNs

$$y = f(x)$$

Generative Model:

Describe the process that generated the data.

Examples:


KDE, GMMs

$$x \sim p(x)$$


Given $X_1, \dots, X_n \sim D$, can we generate X_{n+1}, X_{n+2}, \dots ?

We have encountered a few generative models

$$\begin{aligned} p(\mathbf{x} \mid \mu, \sigma^2) &= \prod_{i=1}^n p(x_i \mid \mu, \sigma^2) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(x_i - \mu)^2}. \end{aligned}$$




$$p(\pi \mid \mathbf{y}) \propto p(\mathbf{y} \mid \pi)p(\pi)$$




$$p(x) = (1 - \pi)\mathcal{N}_{\mu_1, \sigma_1^2}(x) + \pi\mathcal{N}_{\mu_2, \sigma_2^2}(x)$$

$$\ell^t(\theta) \geq \sum_n \left[-D_{\text{KL}}(q_n(z_n) \parallel p(z_n \mid y_n, \theta)) + \log p(y_n \mid \theta) \right]$$



Fitting a probability distribution to real-world data is hard

MNIST



Fashion MNIST



Probabilistic modeling is nevertheless essential

Conditional generative model:

$$p(x|c)$$

Examples:

c = image, x = text

c = initial prompt, x = continuation

c = text prompt, x = image

c = image, x = image



(a) Teddy bears swimming at the Olympics 400m Butterfly event.



(b) A cute corgi lives in a house made out of sushi.

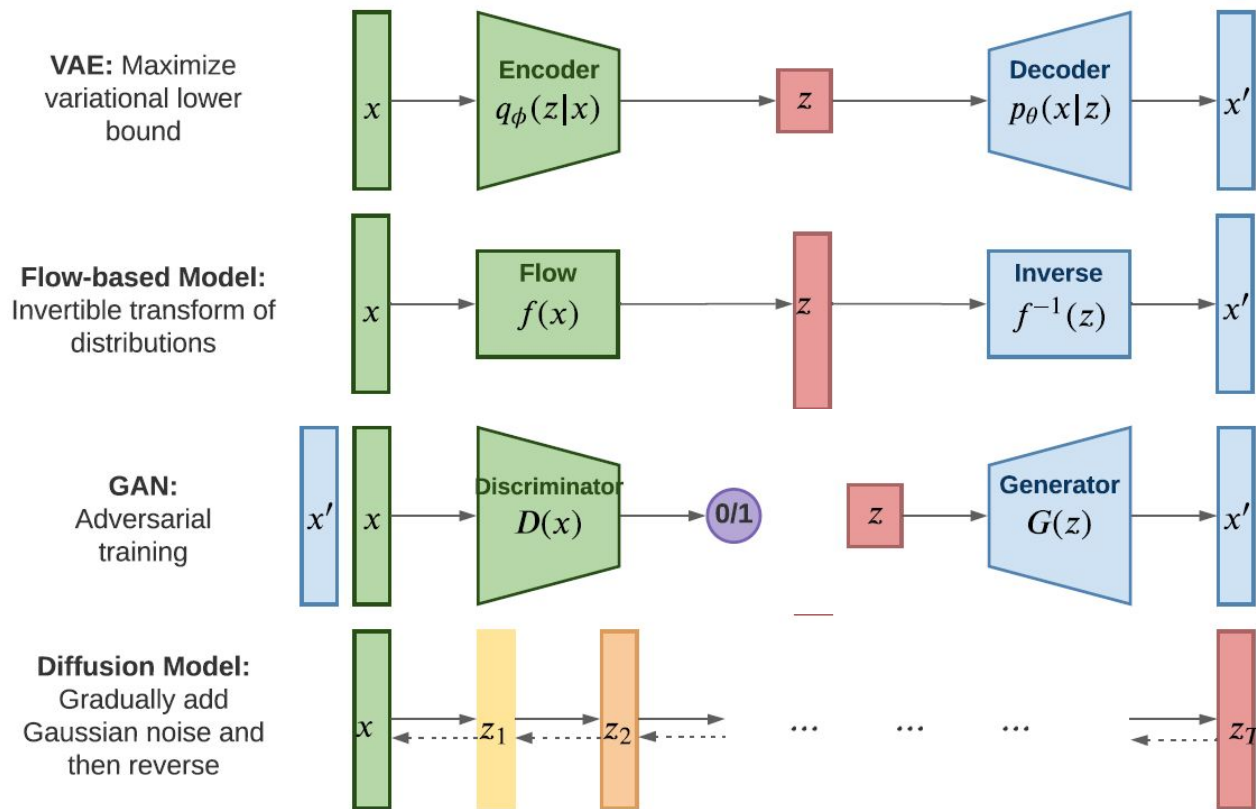


(c) A cute sloth holding a small treasure chest. A bright golden glow is coming from the chest.

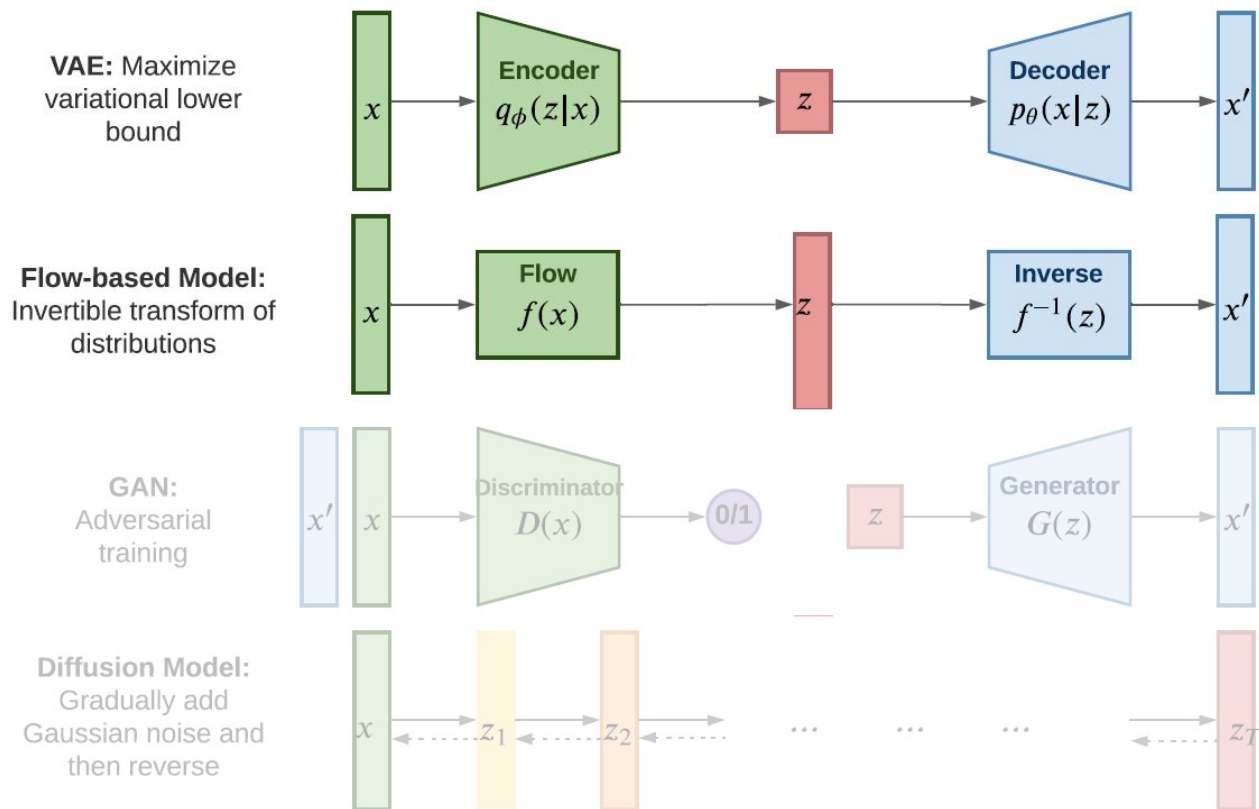
Inpainting




Deep generative models use neural networks to learn a parametrized representation of the distribution



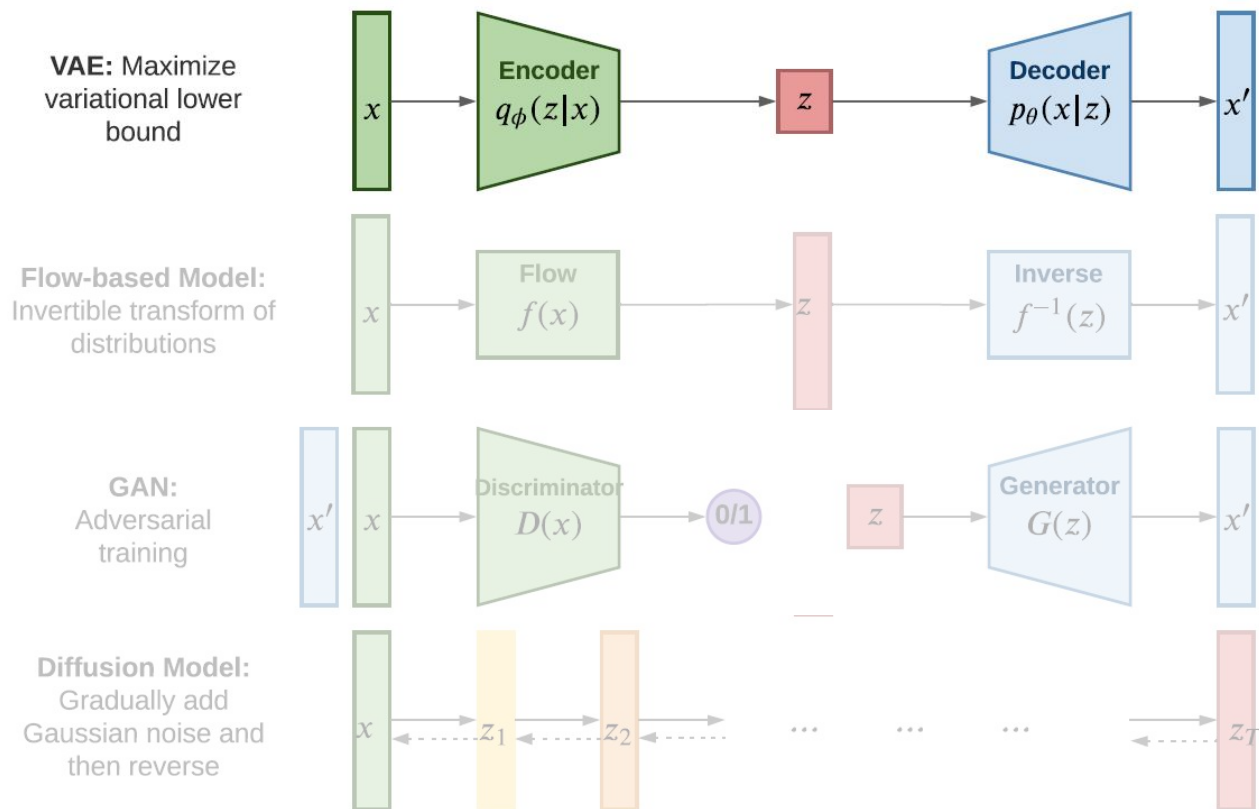
Deep Generative Models: Part I



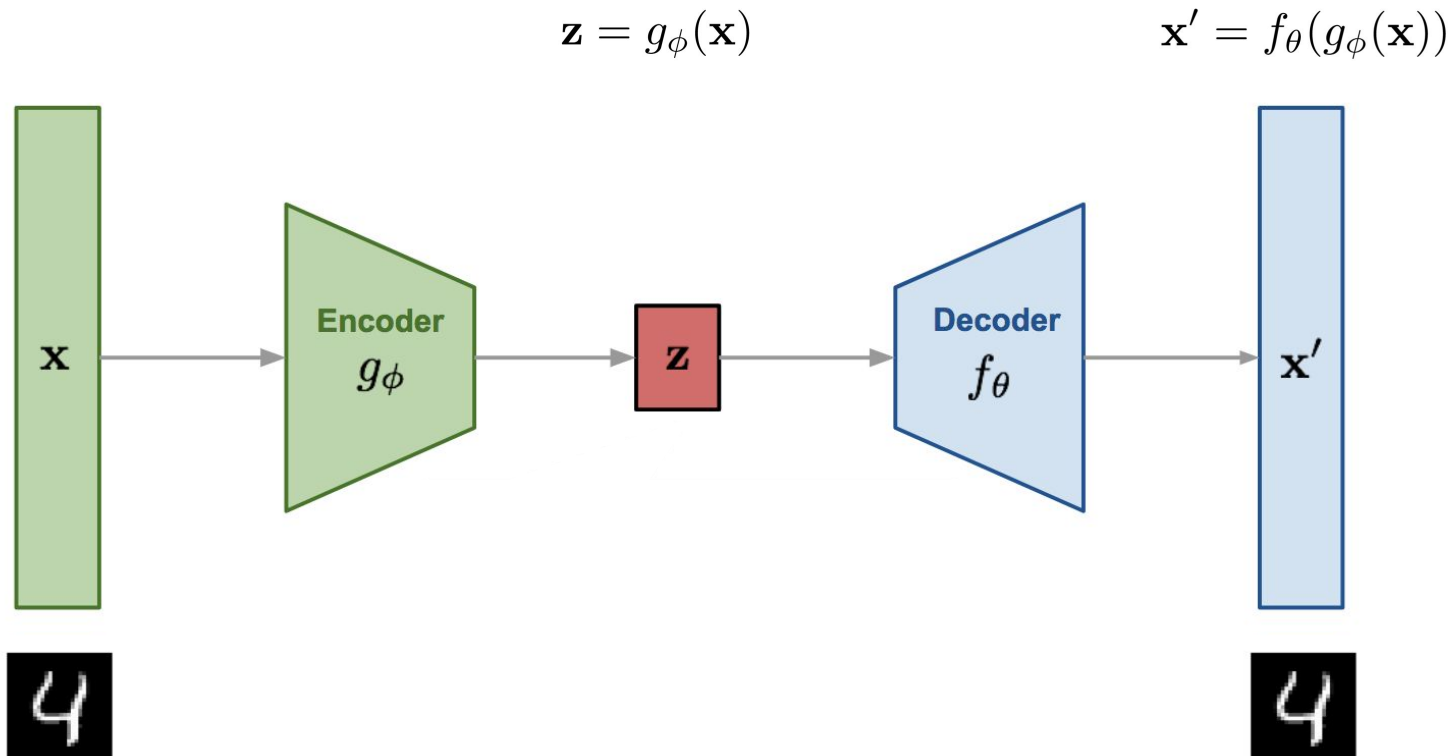
Deep Generative Models: Part II



Variational Autoencoders



An autoencoder is a feedforward neural network trained to learn the identity function



Autoencoders minimize reconstruction loss

Linear:

$$W_f, W_g = \operatorname{argmin} \frac{1}{2} \sum_{i=1}^n \|W_g W_f x_i - x_i\|_2^2$$

Nonlinear:



$$W_f, W_g = \operatorname{argmin} \sum_{i=1}^n \frac{1}{2} \|g(f(x_i)) - x_i\|_2^2$$



Sparse:

$$W_f, W_g = \operatorname{argmin} \sum_{i=1}^n \frac{1}{2} \|g(f(x_i)) - x_i\|_2^2 + \lambda \|f(x_i)\|_1$$



Denoising:

$$W_f, W_g = \operatorname{argmin} \sum_{i=1}^n \frac{1}{2} \|g(f(\tilde{x}_i)) - x_i\|_2^2$$

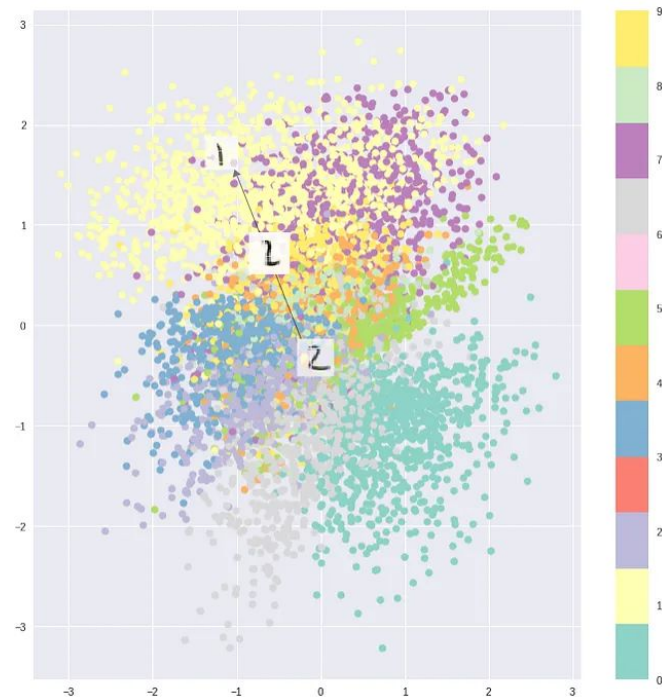





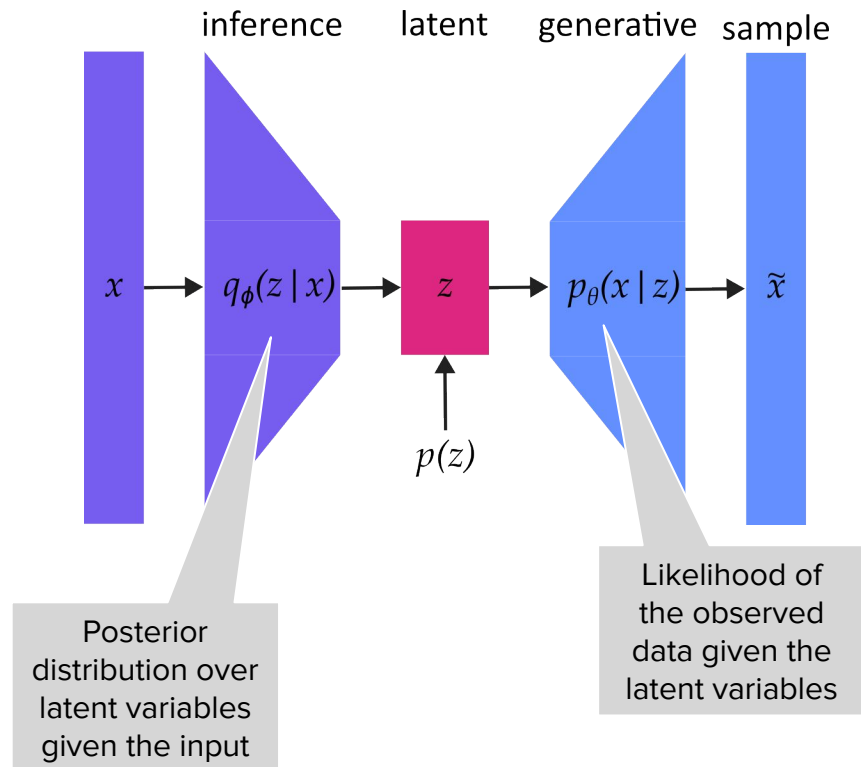

The autoencoder has discontinuities in its latent space



Variational autoencoders encourage continuity in the latent space



VAEs map an input to a distribution



The high-level architecture of a VAE


A VAE defines a generative model of the form

$$p_{\theta}(\mathbf{z}, \mathbf{x}) = p_{\theta}(\mathbf{z})p_{\theta}(\mathbf{x}|\mathbf{z})$$


Where: $p_{\theta}(\mathbf{z})$: is the prior distribution over the latent variable $z \in \mathbb{R}^m$, usually a Gaussian, and $p_{\theta}(\mathbf{x}|\mathbf{z})$: is the density of the decoder network's outputs, conditioned on latent vector.

A VAE approximate the posterior by fitting a “recognition” model:

$$\begin{aligned} q_{\phi}(\mathbf{z}|\mathbf{x}) &= q(\mathbf{z}|e_{\phi}(\mathbf{x})) \\ &\approx p_{\theta}(\mathbf{z}|\mathbf{x}) \end{aligned}$$

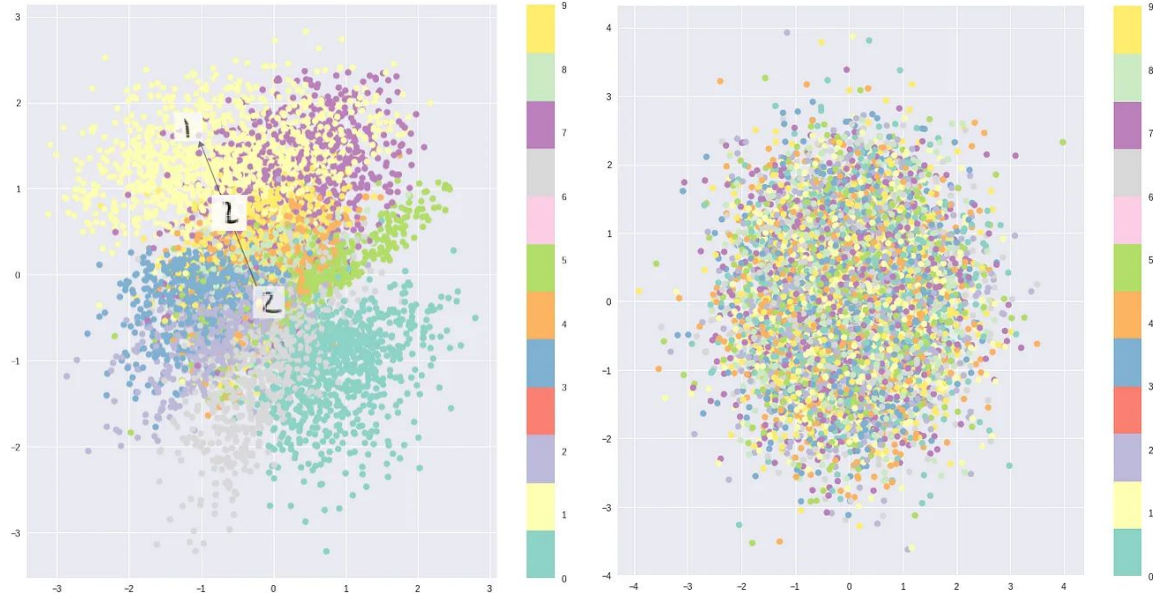


VAEs have both a probabilistic encoder and a probabilistic decoder



The VAE objective balances regularization and reconstruction

0 1 2 3 4 5 6 7 8 9
0 1 2 3 4 5 6 7 8 9
0 1 2 3 4 5 6 7 8 9
0 1 2 3 4 5 6 7 8 9



The VAE objective balances regularization and reconstruction

Encourage the input distribution to correspond to the latent distribution (regularization):

$$\begin{aligned} & \min \text{KL} (q_\phi(\mathbf{z}|\mathbf{x}) \| p_\theta(\mathbf{z})) \\ \implies & \min \text{KL} (q_\phi(\mathbf{z}|\mathbf{x}) \| \mathcal{N}(0, I)) \end{aligned}$$


Encourage accurate decoding (reconstruction):

$$\max \mathbb{E}_{\mathbf{z} \sim q_\phi(\cdot|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})]$$

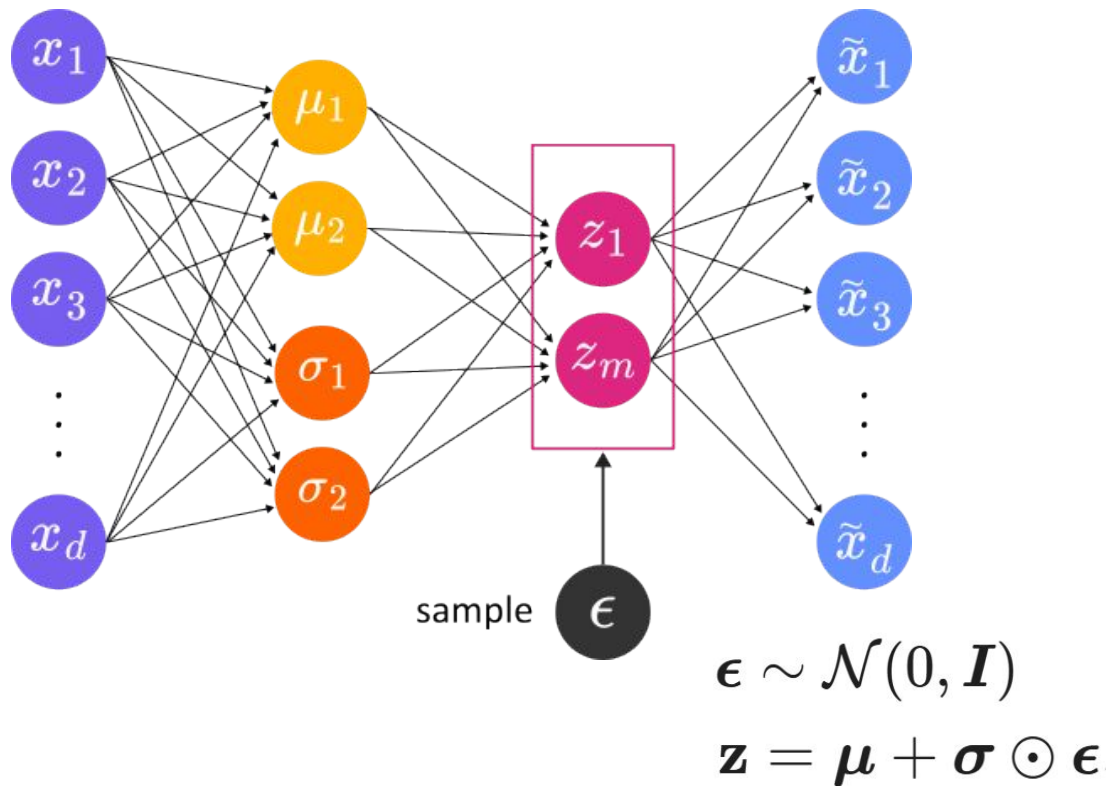
It can be shown that these two quantities define the lower bound on the evidence $p_\theta(\mathbf{x})$:

$$\log p_\theta(\mathbf{x}) \geq \mathbb{E}_{\mathbf{z} \sim q_\phi(\mathbf{z}|\mathbf{x})} [\log p_\theta(\mathbf{x}|\mathbf{z})] - \text{KL} (q_\phi(\mathbf{z}|\mathbf{x}) \| \mathcal{N}(0, I))$$

How to compute gradients across random operations?



The “Reparametrization Trick”



Interpolation with VAEs


[illegible]

AE and VAE on an unconditioned generation task


Training



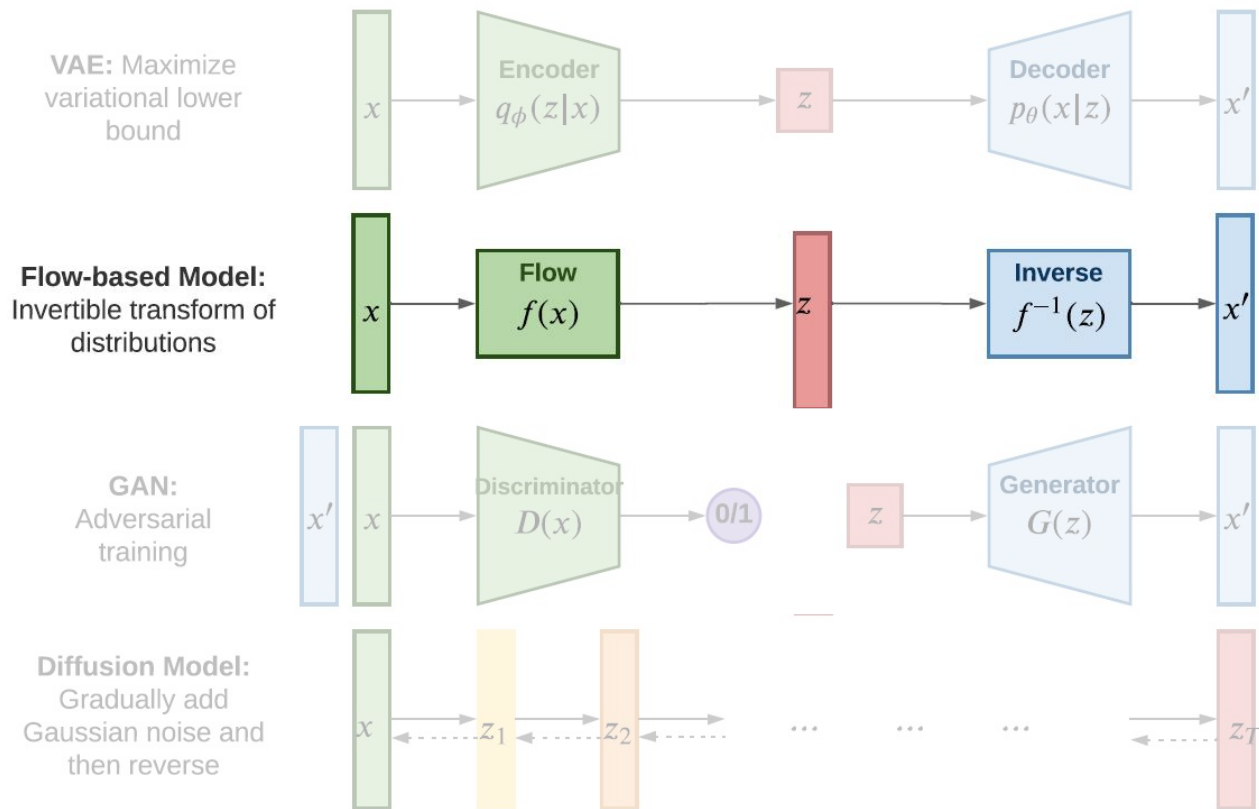
AE
Samples



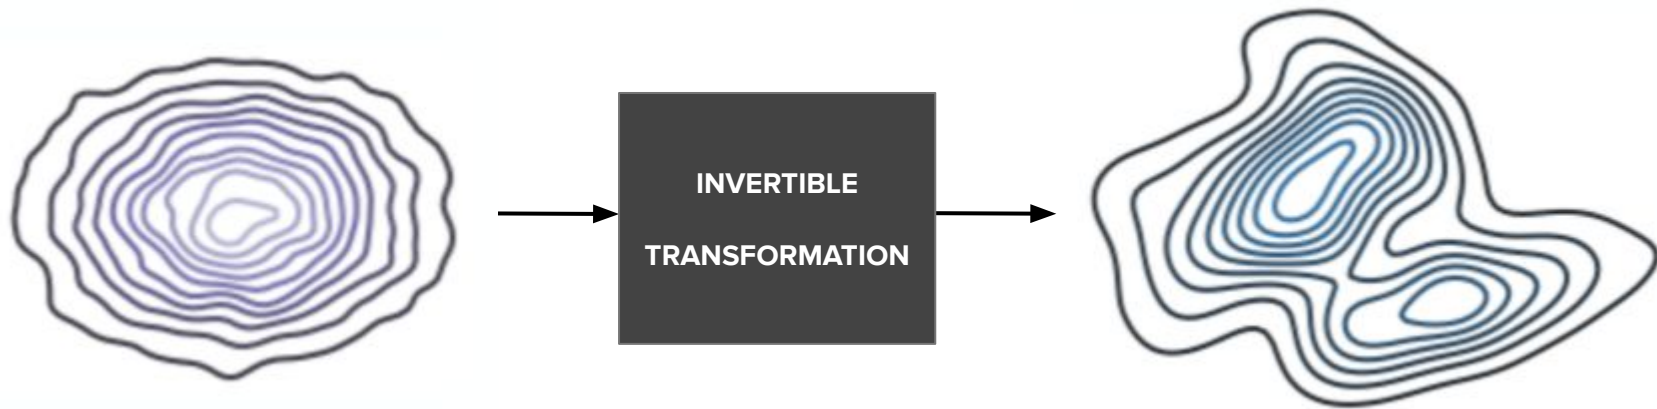
VAE
Samples



Normalizing Flows



Normalizing Flows



Transformations of random variables must preserve probability mass

Uniform random variable:

$$X \sim \text{Uniform}(0, 1)$$

Transformation of X :

$$Y = f(X) = 2X + 1$$


Differentials:

$$(x, x + dx) \rightarrow (y, y + dy).$$


Relationship between probability densities:

$$p(x)dx = p(y)dy$$

$$p(y) = p(x) \left| \frac{dx}{dy} \right|$$



Linear transformations of multivariate distributions must be scaled by the determinant of the projection matrix



$$X = \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix}$$


$$T = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

$$Y = TX$$

$$= \begin{bmatrix} 0 & 0 \\ a & b \\ c & d \\ a+c & b+d \end{bmatrix}$$

$$\int p(y) dy = ad - bc$$


$$= |\det(T)|$$



Smooth nonlinear transformations are locally linear

Given a mapping $f : \mathbb{R}^d \mapsto \mathbb{R}^m$, the Jacobian matrix, defines all first-order partial derivatives:

$$\mathbf{J} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \cdots & \frac{\partial f_1}{\partial x_d} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \cdots & \frac{\partial f_2}{\partial x_d} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \frac{\partial f_m}{\partial x_2} & \cdots & \frac{\partial f_m}{\partial x_d} \end{bmatrix}$$



$$x = f(u), u = f^{-1}(x)$$

$$\begin{aligned} \int_{\mathcal{X}} p_x(x) dx &= \int_{\mathcal{U}} p_u(u) du = 1 \\ \implies p_x(x) &= p_u(u) |\det \mathbf{J}(f)(u)|^{-1} \\ &= p_u(f^{-1}(x)) |\det \mathbf{J}(f^{-1})(x)| \end{aligned}$$

Density estimation with an invertible transformation


Given a dataset $\mathcal{D} = \{x_1, x_2, \dots, x_n\} \sim p_x(x)$, where $p_x(x)$ is some unknown distribution, we wish to learn the density $p_x(x)$.

$$\begin{aligned}\mathcal{L} &= \prod_{i=1}^n p_x(x_i) \\ &= \prod_{i=1}^n p_u(u_i) |\det \mathbf{J}(f)(u_i)|^{-1}\end{aligned}$$

$$\implies \hat{f} = \arg \max_f \prod_{i=1}^n p_u(u_i) |\det \mathbf{J}(f)(u_i)|^{-1}$$

$$\hat{f} = \arg \max_f \sum_{i=1}^n \log p_u(u_i) - \log |\det \mathbf{J}(f)(u_i)|$$

A normalizing flow defines a sequence of *bijectors*



Now that we're at the end of the lecture, you should be able to...

- ★ Distinguish generative from discriminative models, and recall two models that focus on learning an **explicit representation of the distribution: VAEs and normalizing flows**.
- ★ Differentiate autoencoders and variational autoencoders on the basis of the **stochasticity of their outputs** and the **structure of their latent spaces**, and recommend one or the other for particular use-cases.
- ★ Describe how VAEs generate new samples using the **prior distribution in latent space** and the **decoder**.
- ★ Interpret the **loss function of a VAE** with reference to **KL divergence, prior distribution in latent space, evidence lower bound**.
- ★ Defend the use of a **Gaussian prior** in VAE and the **reparametrization trick** for enabling **backpropagation through random operations**.
- ★ List limitations of VAEs and recommend approaches to **improve performance**.
- ★ Use the latent space of a VAE to **interpolate between points in data space**, given the parameters of a small model.
- ★ Differentiate VAEs and normalizing flows on the basis of offering **explicit density evaluation**.
- ★ Use **invertible transformations** to **compute probability density** of a **target distribution** from a **source distribution**.

